TDT4195: Visual Computing Fundamentals

Computer Graphics - Assignment 3

September 25, 2019

Bart van Blokland
Michael Gimle

Department of Computer Science
Norwegian University of Science and Technology (NTNU)

- **Delivery deadline: October 4th, 2019 by 23:59.**

- **This assignment counts towards 6% of your final grade.**

- You can work on your own or in groups of two people.

- Deliver your solution on *Blackboard* before the deadline.

- Use the Gloom project along with all modifications from Graphics Labs 1 and 2 as your starting point.

- Do not include any additional libraries apart from those provided with Gloom.

- Upload your report as a single PDF file.

- Upload your code as a single ZIP file solely containing the *src* and *shaders* directories found in the Gloom project. All modifications must be done to files in these directories. Not following this format may result in a score deduction.

- All tasks must be completed using C++.

- Use only functions present in OpenGL revision 4.0 Core or higher. If possible, version 4.3 or higher is recommended.

- The delivered code is taken into account with the evaluation. Ensure your code is documented and as readable as possible.

Introduction

This is the final Computer Graphics assignment. As with the second one, start off with your code from the previous assignment.

Individual tasks of the previous two assignments have asked you to implement various pieces of functionality necessary to render a 3D scene. That is, a function to create Vertex Array Objects, the implementation of a controllable camera with an appropriate projection, and the corresponding Vertex and Fragment Shaders.

In this assignment, we're going to put all of that to use, and create a scene that's far closer to something you might find in a modern game. For starters, we'll draw some models with far more triangles than the dozen or so you've drawn thus far. Next, we'll implement some very simple lighting, and finally we'll animate the model using a data structure found in every single major game engine out there, known as a "Scene Graph".

The scene consists of lunar terrain and a helicopter [1]. We'll also (quite literally) make this thing fly.

Some of the things we're asking you to do in this assignment will require you to write some more code than what you're used to from the previous ones. We have therefore done the heavy lifting for you, and have provided a handout containing all the code needed to load the models, and for creating and managing the scene graph. We also packaged up the models themselves.

To extract it into your project:

- Put the contents of the "src" directory in the gloom/src directory. Make sure to run cmake again after you do so!

- Place the "resources" directory in the "gloom" folder that also contains the src, shaders, and vendor directories.

Introduction to the Scene Graph

As stated before, we will create and animate a larger scene consisting of a number of different objects. One problem with such scenes is that the transformations for individual objects become increasingly difficult to compute. We will focus on the primary way for modelling scenes, and simplifying transformations within them; the Scene Graph data structure.

---

[1]Yes this is scientifically possible, please stop asking questions.

In its essence, a Scene Graph is a tree-like data structure describing a hierarchy of nodes [2]. This might seem like a somewhat vague definition at first (and you'd be right), but the important bit to note here is the word *hierarchy*.
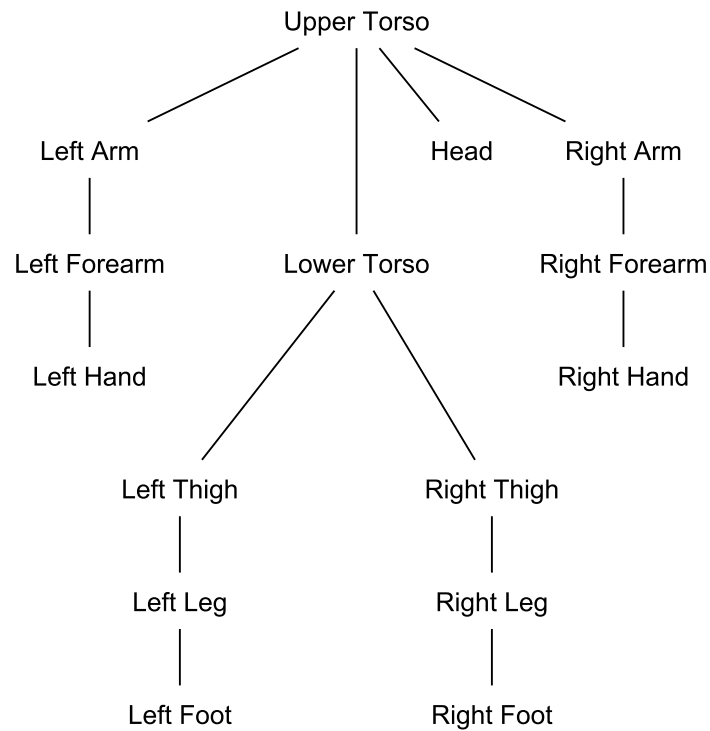
The point behind organising a scene into a hierarchy lies in the fact that it significantly simplifies calculating transformations, and most scenes can intuitively be modelled as such. To illustrate this, let's look at an example.

Consider either one of your arms. When you move it around, notice how your forearm moves along with it. Moreover, since the arm and forearm are connected by a joint, it is impossible to move your forearm independently of your arm (unless something terrible has happened).

In this way, the position of the forearm in 3D space can be described relative to the position of the arm. The same is true for the position of the hand relative to the forearm.

It is possible to model the whole human body in a similar way. Considering the upper torso as a reference, a possible hierarchy is shown in the figure below. Note that toes and fingers have been excluded for simplicity.

---

[2]The best description of the scene graph would be a tree structure: it's a hierarchy of nodes. However, because in some cases it's desirable to reuse child nodes, some nodes can have multiple parents. This means the data structure is technically a graph rather than a tree.

```
                              Upper Torso


          Left Arm                    Head       Right Arm


          Left Forearm          Lower Torso      Right Forearm


          Left Hand                              Right Hand


                    Left Thigh        Right Thigh


                    Left Leg          Right Leg


                    Left Foot         Right Foot
```

From an implementation perspective it is possible to calculate the entire sequence of
transformations of individual parts of a model (such as body parts of a human). However,
it is often far easier to calculate transformations of a particular part of a model relative to
other part(s).

In fact, due to the interdependence of transformations between nodes, a large part of the
transformations you'd apply on for instance a hand are *identical* to the ones on a forearm.
This is precisely what the Scene Graph attempts to exploit.

Additionally, in the case of the movements of your arm, it is far easier to describe how your
forearm moves compared to your arm than how it moves relative to the floor. In terms of
transformations, the main idea of the Scene Graph is therefore that each node *describes
solely how its contents move relative to its parent node.* You can subsequently *combine* the
transformations of each parent node (down from the root node of the tree) with the specific
child node to obtain the complete transformation of that node.

The real power here is that the child node does not need to care what the parent node(s)
are doing. If a character is walking on top of a car, that car can be driving all over the
place, but all you need to do for that character is computing its movements relative to the
car itself.

There are many ways in which nodes in a Scene Graph can be used. Typically, a Scene Graph Node is implemented only as an "interface", which can be implemented by different node types with specific specialisations. It is for instance possible to create nodes which focus on setting values of specific uniform variables, or enable and disable shaders used for rendering particular parts of the scene with specific visual effects.

However, our usage of the Scene Graph in this assignment will focus on on how the Scene Graph hierarchy significantly simplifies the process of calculating transformations of objects. As such we'll only focus on implementing the Scene Graph in its most basic form. This form is equivalent to a tree structure.

First, we'll assume there's only a single type of node. Each node describes its transformation (in a sense movement) relative to its parent in terms of a position, rotation and size (scale). These values will be compiled into a current transformation matrix each frame, also stored in the node.

Next, each node has a Vertex Array Object ID along with the number of indices, which represents the appearance of the node. Finally, each node contains a list of child nodes.

Two of the source files (sceneGraph.hpp and sceneGraph.cpp) which have been attached to this assignment on Blackboard already contain most of the basic functionality needed to implement a scene graph.

Setting up the Scene Graph

Setting up a Scene Graph is essentially equivalent to constructing a tree data structure, as mentioned previously. You create instances of the data structure of each node, and add each child node to its parent's list of children.

You also initialise whichever values need initialisation, such as the initial position and VAO ID of the node (as well as others).

Updating the Scene Graph

Setting up the Scene Graph only needs to be done once. Updating and rendering it is something done each frame. The purpose of updating it is mainly to update positions of animated objects. As with stop-motion animation, the illusion of moving objects can be created by incrementally moving them around a scene many times per second.

It is important here to make a distinction between the two main types of implementing animation: frame-based and time-based. Frame-based animation is by far the easiest to

implement. Each frame, you move an object by a specific amount. Assuming the framerate is constant, this gives perfectly acceptable animations.

Unfortunately, this assumption is also the major downside of frame-based animation. First, the framerate at which a scene can be rendered can slow down significantly as the scene becomes more complex, as each individual frame takes longer to draw. This can also depend on hardware. On the other hand, if the hardware is capable of rendering past a specific framerate for which the animation was designed, animations can appear unnaturally fast.

The solution here is to make the speed of animations depend on a measure which is more consistent; time itself. The idea is to make the displacement of objects depend on how much time has elapsed since the previous frame (a function has been provided for this in the handout code). This means that the displacement of objects is greater each frame the more time has elapsed since the previous frame.

The way you commonly implement this is by once per frame requesting the time that has elapsed since the previous one. You then scale/multiply the movement distance of your object by the elapsed time.

Updating the Scene Graph is a matter of iterating over each node in a depth-first-search order. The only difference is that the parent node is evaluated before its children because child transformations depend on those of its parent, as described previously.

In our case, we'll also determine the correct updated transformation matrix based upon the updated location, scale, and rotation values (see the respective task for detailed instructions).

Rendering the Scene Graph

Rendering a Scene Graph is a process almost identical to updating it. Nodes are traversed in an identical order (compared to updating it), but instead the correct rendering state is set up each time, and a draw call is issued.

Part of this OpenGL state is the complete transformation matrix for the particular node. As the node's transformation matrix depends on the transformations of its parents, its own transformation matrix (the one which describes the transformations relative to its parent), must be multiplied with the one of its parent in the correct order to obtain the accumulated model matrix. It's also necessary to apply the view and projection matrices before sending them to OpenGL.

This is usually accomplished with a matrix stack. The idea is that each node which requires transformations calculates its relative transformation, multiplies it with the one on the top of the stack, and pushes it on to the stack. Note the word "requires" here, because in practice many nodes in the scene graph do not affect transformations. They can for

instance affect render settings or pipeline properties such as lighting coefficients (for phong lighting) or change shaders. Because of that limitation a stack is essential.

Functions have been provided for a stack-based implementation, though passing arguments through recursive function calls may also be intuitive in the case of this assignment too.

Note: you're allowed to make modifications to the supplemental source files, if desired. You may for instance want to add information to each node regarding which shader to use, in case you want to experiment with multiple shaders.

Reference Points

Before we start on the main objective of the assignment, we first should take a detailed look at how movement relative to a parent node can be achieved when using a scene graph.

Specifically, it's worth noting that objects can be said to move relative to their parent (such as your hand in relation to your forearm), but often the points which objects rotate *around* are not the same. Your hand rotates around a joint at the end of your forearm, while your forearm rotates around a joint in your elbow.

For this reason scene nodes in a scene graph tend to use *reference points*, which define the origin of a child node relative to its parent.

To illustrate this, let's focus on a model of a bike, as shown in figure 1. We'll assume we only want to animate the front wheel, back wheel, the pedals, and the steer (so we can turn the front wheel left and right). We'll therefore disregard smaller moving components such as the brakes and chain.

Referencing figure 1, notice that two reference points have been marked; a green one for the bike frame, and a yellow one for the back wheel. Let's assume that all vertices of the bike have been specified relative to the reference point of the bike frame (the green point). Let's also assume we have stored the vertices of each movable part mentioned previously in separate VAOs.

Notice that the bike's origin is placed near the ground. Doing so makes it much easier to tilt the bike up or down when climbing up a hill, or riding down one, which is a single rotation about the z-axis.

The back wheel's reference point has been placed at the centre of the wheel's rotation axis, so that rotating the wheel can be done by (at some point) performing a rotation around the z-axis.

Figure 1: The location of reference points where they could be expected on a 3D model of a bike. Image adapted from: `https://commons.wikimedia.org/wiki/File:Bicycle_diagram-en.svg`

The main problem here is that vertices in 3D models are commonly referenced relative to the overall origin of the model, rather than the each movable object inside of them. For instance, a vertex part of the bike's back wheel would have a negative x-coordinate and a positive y-coordinate, as the origin it's specified relative to is the green reference point.

This phenomenon is very common because objects are often created or edited in 3D modelling software. These tools will produce models whose coordinates are all specified relative to the same origin.

Rotating the back wheel around the reference point of the 3D model yields the effect shown in figure 2. This is generally not how bike wheels behave when rotating.

In order to rotate the back wheel around the reference point, you would first have to move it to the origin, then apply the rotation, and finally move it back to where it was. Conveniently, a movement to the origin is accomplished simply by translating by a vector which is the inverse of the reference point.

Figure 2: The effect of rotating the back wheel around the origin of the 3D model. The red arrow indicates the rotation motion. Image adapted from: `https://commons.wikimedia.org/wiki/File:Bicycle_diagram-en.svg`

Task 1: More polygons than you can shake a stick at [1 point]

To start out this assignment, we will use the code you wrote in the last assignment to draw a relatively detailed lunar surface model, consisting of close to 100 000 triangles.

You see, modern games don't draw a triangle or two at a time, they draw anything between tens of thousands to millions of triangles. And graphics cards are nowadays powerful enough to be able to do that. So for this first task, you'll get to see what it takes to go from the 5 triangles you've drawn, to a few hundred thousand.

And here's the real kicker: the major difference between drawing a few triangles rather than hundreds of thousands is... Absolutely Nothing!

It turns out that what you've already done in the previous assignment is exactly what you need to do in order to draw as many triangles as you want!

Let's dive in, shall we?

a) **[0.4 points]** We'll start out by loading the terrain model.

Since this model covers a fair amount of distance (the difference between the left-most and right-most coordinates is relatively large), you probably want to increase the *far plane* of your perspective matrix before you start. You could for example use $1000.0f$.

Up to this point we've been defining individual triangles by manually specifying a vertex and index buffer. In practice, you usually never do this, perhaps apart from specifying shapes of a few triangles at most.

If you look through the handout zip file, you will find a folder named "resources". This contains two `.obj` files, one for a helicopter (which we'll be using later), and the other for a piece of terrain. Place the folder in the `gloom` folder of your project, next to `shaders` and `src`, so you can access it in a similar way to how you accessed the shaders in the previous assignments.

Load the `lunarsurface.obj` mesh and create a VAO from it, and then draw the VAO. Make sure your call to `glDrawElements()` draws the right number of indices!

You won't have to figure out how you can load this model yourself. We've done this for you, and included all the data structures and functions you need in `OBJLoader.hpp`:

```
Mesh loadTerrainMesh(std::string const srcFile);
```

A `Mesh` contains the following variables:

```cpp
struct Mesh {
    std::string name;
    std::vector<float> vertices;
    std::vector<float> colours;
    std::vector<float> normals;
    std::vector<unsigned int> indices;
};
```

If you want to use another terrain model, you're free to do so, but keep in mind that the model loader that comes with the code handout only supports `.obj`.

b) [**0.2 points**] The model is going to look pretty boring when coloured with a solid colour. We want to see some of the juicy details present in the decently high triangle count model. In order to accomplish this, we are going to be adding some very simple lighting.

In general, the primary factor determining how much light is reflected from a surface, is the angle between the light source (the sun, a light bulb, etc), and surface itself.

So if we have to compute the colour of a given fragment, we first have to know what *angle* it makes with the light source. In order to do this, we have to know what direction the triangle is facing. This is usually achieved by the so/called "normal" vector. This is a unit vector (of length 1) that points towards the "outwards" side of the surface.

Because we know which direction our triangle is facing, we can of course easily compute this. However, the more common way of solving this is to define a single normal for each vertex in your model. As you may have seen in the Mesh definition in the previous task, the object you load in already contains a list of normals that the .obj file defined. Similar to the colour buffer in the previous assignment, we now have to make these normals a part of our VAO.

At this end, extend your VAO generation function to also take in a vector of floats containing the `x`, `y` and `z` coordinates of the normal vectors included in the mesh. This is basically doing exactly the same as what you did in the second assignment.

c) [**0.2 points**] [**report**] Extend the *vertex* shader to take the new normals as an input, and pass them straight on to the fragment shader through an output variable.

Extend the *fragment* shader to take in the normals that the vertex shader outputs.

After you've done this, visualize the normals for a model by using the `x`, `y` and `z` coordinates of the normals as the `r`, `g` and `b` values for our fragments.

This will make some of the colours negative (i.e. black), but that's fine for this example.

When you do this, your scene should end up looking very colourful. Lots of shades of green, red, and blue. Usually no grays, though.

Position your camera to point into one of the craters in the scene, and attach a screenshot of this immense natural beauty in your report.

d) **[0.2 points] [report]** Finally, we want you to implement some very simple lighting.

In the fragment shader, create a variable holding the direction that the light is coming from. We're creating a light source that's infinitely large here, which on earth effectively is the sun.

```
vec3 lightDirection = normalize(vec3(0.8, -0.5, 0.6));
```

Computing light accurately is immensely complicated and computationally intensive. So instead we're going to use what's known as a "lighting model"; an often crude approximation that looks plausible to the human eye.

The lighting model we are going to be using assumes that every object reflects light equally much in all directions (obviously not the case, but it's easy to compute). This is called a Lambertian model. You can read more about it at `https://en.wikipedia.org/wiki/Lambert%27s_cosine_law`, but it's not required for this assignment.

The equation we want you to use to compute the final colour in the fragment shader looks like this:

$$colour * max(0, normal \cdot -lightdirection)$$

When you set the fragment colour to this, all of the details in the lunar surface should be visible.

Attach a screenshot which shows that the surface is correctly lit.

Task 2: Helicopter Parenting [1.0 point]

Now that we have a piece of terrain, we can move on to something that moves around in it. Or rather, *flies* around in it. We'll therefore now be loading the helicopter model that comes with the handout archive, and draw it.

Also, we'll be constructing the Scene Graph data structure which we talked about before.

a) **[0.2 points]** First of all, load the helicopter model, found in `helicopter.obj`, using the provided function:

```
Helicopter loadHelicopterModel(std::string const srcFile);
```

A `Helicopter` contains the following:

```
struct Helicopter {
    Mesh body;
    Mesh mainRotor;
    Mesh tailRotor;
    Mesh door;
};
```

Create a VAO for every part of the helicopter. Draw the VAOs for the helicopter model.

b) **[0.5 points]** Create a function which constructs and returns a scene graph containing:

- A node set up to render your terrain.
- At least one node containing and instance of the helicopter you made in the previous task.

To accomplish this task, do the following:

1) Generate one SceneNode for each object (Terrain and helicopter parts) in the scene.
2) Organise the objects into a Scene Graph by adding child nodes to their parent's list of children. The organisation must be logical in terms of which object(s) should move relative to other objects, as described previously.
3) Initialise the values in the SceneNode data structure to their respective initial values, such as the position and starting rotations.
4) Connect the root node of the helicopter to the terrain, and the terrain to a single root node for the entire scene. Return this node from your function.

Relevant functions in the sceneGraph.hpp file:

- `SceneNode* createSceneNode()`
- `void addChild(SceneNode* parent, SceneNode* child)`
- `void printNode(SceneNode* node)`

c) [**0.3 points**] [**report**] Now that we have the scene structured into a scene graph hierarchy, we can use it to determine what to draw instead of just calling the draw function for each VAO manually.

We won't be needing the code you wrote to draw the VAOs in the earlier tasks, so you can now get rid of the drawing code in your main loop.

Drawing the scene involves iterating over every `SceneNode`, binding its VAO and then calling `glDrawElements()`.

This snippet should get you most of the way there:

```
void drawSceneNode(SceneNode* root, glm::mat4 viewProjectionMatrix) {
    // Draw here

    for (SceneNode* child : root->children) {
        drawSceneNode(child, viewProjectionMatrix);
    }
}
```

Note the viewProjectionMatrix parameter. You'll need to use it later, but don't worry about it for now. Just pass in the View Projection matrix you're computing in the main loop, and call it a day.

Attach a screenshot showing the helicopter being drawn.

Task 3: The (Model) Matrix: Revolutions [1.5 point]

We've loaded in some models now, and they're looking very pretty with their hundreds of thousands of triangles. But there's a problem here: our helicopter is not moving. How do we go about moving things around the scene?

So far we've seen the projection and view matrices, and how they combine to allow us to see the scene from any viewpoint (as we did in the previous assignment). The third matrix commonly used in rendering is called the "model matrix", and represents transformations of objects/models (or parts thereof) *within* the scene. The model matrix is usually applied before the view and projection matrix, hence the name by which the combination of these three matrices are commonly referred to; The Model View Projection (MVP) matrix.

a) **[0.3 point]** The first thing we need to do is to implement a function which traverses the Scene Graph. It should be able to "visit" each node (use a depth first search or similar). We'll implement what to do at each node later.

The following snippet should get you mostly there:

```
void updateSceneNode(SceneNode* node, glm::mat4 transformationThusFar) {

    // Do transformation computations here

    // Store matrix in the node's currentTransformationMatrix here

    for(SceneNode* child : node->children) {
        visitSceneNode(child, combinedTransformation);
    }
}
```

b) **[0.2 points]** In the SceneNode struct (see Scenegraph.hpp), there's a referencePoint field. This corresponds to the point about which the contents of the node should be rotated. This point should be defined in the coordinate space of the object contained within the node's VAO.

That basically means you look at the object in the VAO stored in the node, and choose a point which that part should be rotated around.

Correctly set the reference point of all nodes in your scene graph.

To avoid a wild goose chase, the tail rotor reference point of the helicopter model is at $[0.35, 2.3, 10.4]$.

c) **[0.7 point]** Another field in the SceneNode struct is the `currentTransformationMatrix` one. This matrix represents the node's transformation relative to its parent. Each time the traversal function you created earlier visits a node, compute an updated version of the node's matrix based on the node's location (x, y, z) and rotation (rotationX, rotationY, rotationZ) fields.

The node should be rotated around its reference point.uires some transformations in addition to the rotations themselves. It may be useful to refer back to the introduction of this assignment for this.

d) **[0.3 point]** The one thing remaining is that when we want to draw a specific node, we need to combine the node's Model matrix which we just computed with the scene's View Projection matrix.

Therefore, in the `drawSceneNode` function, combine the View Projection matrix (that you should already be passing around at this point) with the node's `currentTransformationMatrix`. Then, pass it into the vertex shader by setting the

value of this transformation matrix uniform you've created in the previous assignment to this new matrix. Again, make sure the multiplication order is correct here!

If you want to verify that it works, you can change the position or rotation values in the `SceneNode` of the helicopter body.

Task 4: Spinning into gear [1 point]

Time for some animations! Now that you have set up the scene graph properly, this should be a cinch. All animations should be done in a time-based rather than frame-based way. This is easy if you use the function we provide in `toolbox.hpp` called `getTimeDeltaSeconds()`. This function returns the amount of time in seconds that has elapsed since the last time it was called. Therefore, make sure to only call it once per frame.

a) **[0.4 points]** Make the helicopter's main rotor and tail rotor spin in their sockets continuously.

   The order in which you multiplied the matrices in the previous task when updating the scene graph is very important. If you're not sure if you have done it right, try working through the next subtask, as it should look very obviously wrong if your order is incorrect.

b) **[0.6 points]** For this task, you should have your helicopter follow some animated path. We have provided a function which will generate some coordinates and rotations following a circuit.

   You're free to come up with any animation you'd like here, but it's important that you showcase rotation along all three axes if you do. This animation is going to look wrong unless you have correctly implemented all parts of the Scene Graph update.

   Use the following function to get all the information you need to animate the helicopter.
   `Heading simpleHeadingAnimation(double time);`

   A `Heading` contains the following:

```
struct Heading {
    float x;
    float z;
    float yaw;
    float pitch;
    float roll;
};
```

Task 5: Help! My lighting is wrong! [1 point]

a) **[0.2 points]** **[report]** Notice anything weird about how the helicopter is lit as it moves around? The lighting doesn't change even though the helicopter is turning!

This is because the normal vectors are defined in relation the model as if the model still had its original orientation. If we want the shadow to move across the helicopter when the helicopter rotates, we have to rotate the normals in the same way that we rotate the model.

Attach two screenshots taken from the same camera position, one where we see the brightly lit side of the helicopter, and another taken when it has rotated and we see a darker side instead.

b) **[0.4 points]** In the real world, solving this problem requires the use of something called a Normal Matrix, but since that's outside of the scope of this course, we'll use a cheap imitation instead. In short, the idea is that we butcher our existing model matrix to compute the transformation matrix that rotates the normals to the correct orientation.

Note that we need to rotate normals using the Model matrix only; we don't want a matrix that also has the View and Projection matrices applied on it.

Fortunately, we have been passing around the View Projection in our traversal function that draws the scene graph, and we have the model matrices of all nodes stored in the SceneNode's currentTransformationMatrix field.

As such, we can compute two separate matrices, and pass each of them in their own uniform variable into the vertex shader:

- The Model View Projection matrix, used to transform the input vertex
- The Model Matrix, used to transform the vertex normal

c) **[0.4 points]** Before we multiply the Model matrix with the vertex normal, we'll need to make some minor modifications to it to ensure it does not scale or translate the normal.

We do this by taking the top 3x3 part of the Model matrix, we get the all of the rotation applied to the model, and none of the translation. We then multiply this by the normal vector, which results in it being rotated by the same amount that the object it's attached to is. Finally, we normalize the result, since any change in the normals tends to result in them getting a magnitude not equal to 1.

```
vnormal = normalize(mat3(Model) * normal);
```

Hint: Don't forget that you can always visualize the normals like we did in the previous task to get a more clear picture of what's going on!

Task 6: Time to turn this thing up to ~~11~~ 5 [0.5 point]

a) [**0.5 points**] Everything you have done so far should be easily extendable to be done in loops, so that's what we're going to do for the final part of this assignment.

Instead of creating a single helicopter when reading the model, use the same meshes to instantiate at least **5** helicopters. Keep track of the pointers to the relevant scene nodes and animate them in the main loop. Every helicopter should have a rotating tail rotor and main rotor, as well as follow some kind of path like in the previous task.

Every one of the helicopters should follow the same path, but none of them should collide with any other. You can accomplish this by feeding the animation function a carefully selected offset.

Task 7: Optional Challenges [At most 0.51 points]

This task is optional. These questions are meant as further challenges or to highlight things you may find interesting. They can reward up to 0.51 points, to supplement missing points from other parts of the assignment. Please show us some eye candy if you manage to answer some or all of them!

a)   Implement Phong shading instead of Lambertian shading [**0.2 points**]

The Phong lighting model is a more visually interesting and more realistic alternative to what we have used in this assignment. It's comprised of three parts, where the lambertian diffuse shading model we used is one of them. If you want a good challenge with coordinate spaces, shaders, and vectors, I recommend giving this a try.

b)   Make one of the helicopters controllable [**0.1 points**]

It should be controllable in the same way the camera was in assignment 2 if one did the bonus task, so going forward should move the helicopter in the direction it's facing. Having the helicopter tilt somewhat realistically is recommended, but not required. If doing this task, we suggest you either leave the camera staying still, or implement the chase camera below (other advanced camera control functionality can also be rewarded).

c)   Implement a chase camera [**0.2 points**]

A chase camera has a position, a target and a chase radius.

It will constantly look at its target, and if the target leaves the chase radius, it will move towards the target in such a way that the target (barely) reenters the chase radius. Changing the animation function for this task is not required, but could be a good idea in order to showcase the efficacy of your chase camera.

d)   Make the door animated [**0.1 points**]

The door of the helicopter is a separate object. This means that it should be pretty easy to animate. Make one of the keys on the keyboard open the door.

You can simply slide it in the Z-direction, but we're very open to see some creative solutions here!

Closing the door again is optional.

e)   Find the easter egg [**0.01 points**]

Attach a screenshot of the easter egg we have hidden somewhere in the scene.

f)   Impress us! [**0.0-0.5 points**]

This is a wildcard for us to give you bonus points if you do anything that's really cool or beautiful. These points are not cheap, though.